

Joseph Carluccio
BisonBnB Consulting
Lewisburg, PA 17837
April 2024

Dear Alia Stanciu:

On behalf of BisonBnB Consulting, we thank you for this opportunity. The goal was to accurately predict the pricing of an Airbnb Listing in New York City, using property and listing attributes. As you know, this prediction can be crucial for hosts and guests, as it aids hosts in setting competitive prices and assists guests in making informed booking decisions.

Given the provided data, we have completed an exploratory data analysis prepped the data, and then built and tested the performance of 4 different prediction models. Based upon our analysis, we would recommend implementing the gradient boosting model to most accurately predict the listing price as this one had the best performance based on various metrics.

DATA: SUMMARY AND GENERAL PREPARATION

The dataset consisted of information about 48,895 Airbnb listings across the 5 boroughs of New York City. The information about each listing was 16 different attributes - 11 quantitative and 5 categorical.

The information about each listing:

listing ID

name - name of the listing

host ID

host_name - name of the host

neighbourhood_group - which of the 5 boroughs is the listing located in

neighbourhood - neighborhood the listing is located in

latitude - latitude coordinates

longitude - longitude coordinates

room_type - listing space type

price - price in dollars

minimum_nights - the amount of nights minimum

number_of_reviews - number of reviews

last_review - date the last review was given

reviews_per_month - number of reviews per month

calculated_host_listings_count - the number of listings per host

availability_365 - number of days when the listing is available for booking

Missing Data:

Of the data provided, 10052 listings had at least one null data field.

Of the fields missing, they were all in fields relating to reviews or the name of the host. So, we filled these missing values with 0 as they were not important to our Regression Models.

We addressed these missing values before building any of the models. We thought this would give the best comparison of performance across them.

Outliers:



Figure 1.

Scatterplot of Airbnb Prices

Of the 48,895 listings, 6021 were identified as outliers. These outliers can be seen in Figure 1 with some Airbnbs being priced as high as \$10,000. Since this is quite a significant portion of our dataset, these outliers were left in the data. Also, since 2 of our implemented models are

tree-based (which are naturally less likely to be affected by outliers), these outliers should have very minimal effects on the quality of the analysis and the performance of these models.

Training and Validation Data:

For each of the 4 models, the data was split into 80% training data (39116 listings) and 20% testing or validation data (9779 listings)

MODELING:

4 different models were implemented all with the same parameters and target value

After the initial EDA, we determined which parameters had the most effect on price: neighborhood group, latitude, longitude, room type, minimum nights, number of reviews, last review, and availability 365. The target value for all of these models was price.

Selected Models

1. Decision Tree Regression:

Performed Grid Search to tune max_depth, min_samples_split, and min_samples_leaf.
Found that the optimal values for each were:

max depth, which controls the maximum number of levels in the tree = 7

minimal samples per leaf, which controls the minimum number of samples required to consider splitting a node further = 10

minimal samples per split, which controls the minimum number of samples required to consider a node as a leaf (i.e., not split further) = 2

Tested the performance of the model using 5-fold cross-validation with negative mean absolute error as the scoring criterion. The output showing the Decision Tree can be seen in Figure 1.

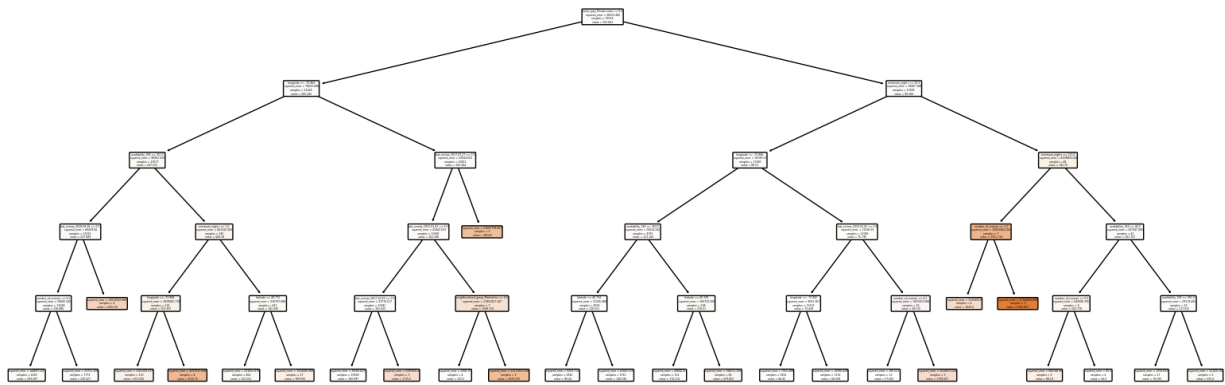


Figure 2.

Decision Tree Regression Model

2. LassoCV (Lasso Regression with Cross-Validation):

Implemented using basic sci-kit learn functionality. To be used in comparing Decision Tree Regressor performance. Grid Search CV was not needed for LassoCV, as it handles hyperparameter tuning internally

3. Gradient Boosting Regression:




Figure 3.

Scatterplot Showing the Actual Vs. Predicted Price using Gradient Boosting Model

The Gradient Boosting Model is similar to a Decision Tree Regressor. By fitting smaller decision trees to the residuals of the previous trees, gradient boosting can gradually improve the predictive performance of the model.

We attempted to perform hyperparameter tuning on this model but ultimately ran into extensive run times with no foreseeable outcome. We attempted to first perform a Grid Search with all the possible hyperparameters of the model but found this to run for an extended and unknown amount of time. We shortened the possible range of hyperparameters while now only focusing on the maximum depth of each tree and the minimum samples to split the leaf further. This once again resulted in extensive run times with no end. We then tried to use a random search with 100 iterations followed by only 10 iterations. Both of these processes ran for a total of 1.5 hours. perhaps this was due to an error somewhere in the code or to a lack of computing power.

Since we are unable to see the possible hundreds of trees that the model creates, we wanted to visualize the performance of the model using Figure 3. As we can see, this model does a good job of predicting the price of the listings, even with some of the outliers.

4. Dummy Model:

The Dummy model predicts the price of the Airbnb Listing by simply taking the average price of all the Airbnb listings included in the dataset. This is heavily skewed by any outliers and clearly does not take into account any characteristics of the listings.

MODEL PERFORMANCE

In the context of our goal,

MAE is the average value that our model's prediction is off from the actual price. This can be in either the negative or positive direction.

MSE is the average difference between predicted and actual prices, squared.

RMSE is the square root of the MSE. This metric gives us a better idea of the actual performance of the model as it minimizes the effects of outliers.

Performance of Models

Model	DT Regression	LassoCV	Gradient Boosting	Dummy (mean)
MAE	69.73	74.92	69.02	93.34
MSE	49996.24	49722.40	50282.01	55028.32
RMSE	223.60	222.24	224.24	234.60

Table 1.

Performance of Models

From Table 1, we can see that the best-performing model was the gradient boosting model with Our decision tree regression performing second best, LassoCV performing third best, and our dummy model performing the worst.

Our best model was the Gradient boosting model with a prediction that was on average, off by \$69.73 of the actual price.

This was a significant improvement over our worst model, the dummy model which on average, was off by \$93.354 of the actual price.

IMPLEMENTED MODEL FOR PERFORMANCE

From Table 1, we can see that the Gradient Boosting model performed the best, and as such, that is the one that we would recommend implementing. Our Gradient Boosting model can accurately predict the nightly rate of the Airbnb Listing within \$70. obviously, in practice, this is not a good tool to give you an exact number to include in the listing, but does offer value for the host to have an estimate of the value of their listing, and even more value for the guest as they can examine the listing and judge whether it is worth the extra \$70 (or \$70 less!).

One thing that we would suggest to increase the performance of the model is hyperparameter tuning to further optimize the model. We attempted to do so with a Grid Search and Random Search but ultimately ran into a limit with computing power. In the near future, we are upgrading our computers here at BisonBnB and would love the opportunity to continue working on the model.

SECTION 5. APPENDIX

```
In [ ]: # import packages
import warnings
warnings.filterwarnings('ignore')

import os

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import randint

import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.linear_model import LassoCV
from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.dummy import DummyRegressor

# Set the display options to print the dataframe
# pd.options.display.float_format = '{:,.3f}'.format
```

We are going to try and predict air bnb prices in NY based off of information about the airbnb

```
In [ ]: # Readin the data
air_df = pd.read_csv('Data/AB_NYC_2019.csv')
air_df.head()
```

Out []:	id	name	host_id	host_name	neighbourhood_group	neighbourhood
0	2539	Clean & quiet apt home by the park	2787	John	Brooklyn	Kensington
1	2595	Skylit Midtown Castle	2845	Jennifer	Manhattan	Midtown
2	3647	THE VILLAGE OF HARLEM....NEW YORK !	4632	Elisabeth	Manhattan	Harlem
3	3831	Cozy Entire Floor of Brownstone	4869	LisaRoxanne	Brooklyn	Clinton Hill
4	5022	Entire Apt: Spacious Studio/Loft by central park	7192	Laura	Manhattan	East Harlem

```
In [ ]: air_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    48895 non-null  int64
1   name                                48879 non-null  object
2   host_id                             48895 non-null  int64
3   host_name                           48874 non-null  object
4   neighbourhood_group                  48895 non-null  object
5   neighbourhood                        48895 non-null  object
6   latitude                            48895 non-null  float64
7   longitude                            48895 non-null  float64
8   room_type                           48895 non-null  object
9   price                               48895 non-null  int64
10  minimum_nights                       48895 non-null  int64
11  number_of_reviews                    48895 non-null  int64
12  last_review                          38843 non-null  object
13  reviews_per_month                    38843 non-null  float64
14  calculated_host_listings_count       48895 non-null  int64
15  availability_365                     48895 non-null  int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB
```

After the quick look at all of our data's features and their types, lets now examine the distribution for our target value, price:

```
In [ ]: # Plot Histogram of the price distirbution

# Make the figure
fig, ax = plt.subplots(figsize = (10,5))
```




```

# set labels
ax.set_xlabel("Price per Night ($)")
ax.set_ylabel("Frequency")
ax.set_title("Price Distribution")

# change x axis to only show values between 0 and 3000
plt.xlim(0, 3000)

plt.hist(air_df['price'], bins = 500, color = 'skyblue')
plt.show()

```



Find the minimum, maximum, median, and mean price in our data

```

In [ ]: # we can get most of our required values from the .describe() function in pa
price_summary = air_df['price'].describe()

# Extracting median separately
median_price = air_df['price'].median()

print("Summary statistics for price:")
print(price_summary)
print("Median price:", median_price)

```

Summary statistics for price:

```

count    48895.000000
mean      152.720687
std       240.154170
min         0.000000
25%        69.000000
50%       106.000000
75%       175.000000
max     10000.000000
Name: price, dtype: float64
Median price: 106.0

```

Now let's examine the distribution for other features of the data. Let's try to focus on ones that will help us best predict the value of price. From real world experience, my assumptions for two features that will give us the most insight are neighborhood, minimum_nights, and room_type.

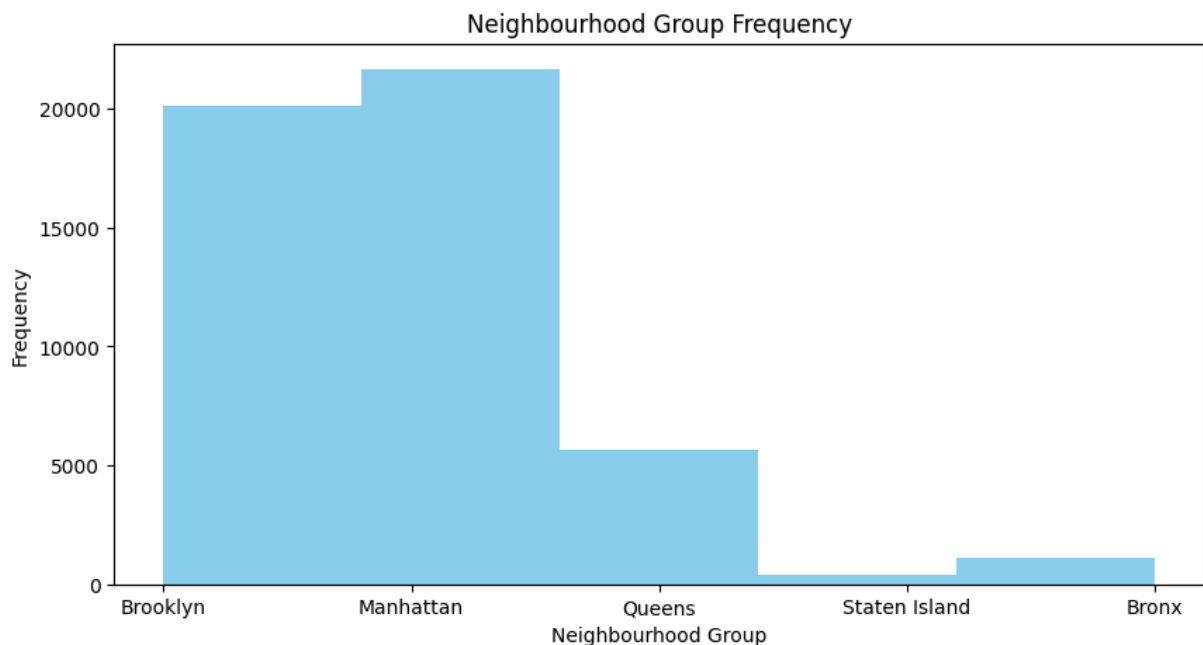
```
In [ ]: # Plot Histogram of the neighbourhood_group

# Make the figure
fig, ax = plt.subplots(figsize = (10,5))

# set labels
ax.set_xlabel("Neighbourhood Group")
ax.set_ylabel("Frequency")
ax.set_title("Neighbourhood Group Frequency")

# change x axis to only show values between 0 and 3000
#plt.xlim(0, 3000)

plt.hist(air_df['neighbourhood_group'], bins = 5, color = 'skyblue')
plt.show()
```




```
In [ ]: # Plot Histogram of the minimum_nights

# Make the figure
fig, ax = plt.subplots(figsize = (10,5))

# set labels
ax.set_xlabel("Minimum Nights")
ax.set_ylabel("Frequency")
ax.set_title("Minimum Nights Distribution")

# change x axis to only show values between 0 and 3000
plt.xlim(0, 400)
```

```
plt.hist(air_df['minimum_nights'], bins = 300, color = 'skyblue')
plt.show()
```




```
In [ ]: # Plot Histogram of the room_type

# Make the figure
fig, ax = plt.subplots(figsize = (10,5))

# set labels
ax.set_xlabel("Room Type")
ax.set_ylabel("Frequency")
ax.set_title("Room Type Distribution")

# change x axis to only show values between 0 and 3000
#plt.xlim(0, 400)

plt.hist(air_df['room_type'], bins = 3, color = 'skyblue')
plt.show()
```




```
In [ ]: # Plot Histogram of the availability_365

# Make the figure
fig, ax = plt.subplots(figsize = (10,5))

# set labels
ax.set_xlabel("Availability 365")
ax.set_ylabel("Frequency")
ax.set_title("Availability 365 Distribution")

# change x axis to only show values between 0 and 3000
#plt.xlim(0, 400)

plt.hist(air_df['availability_365'], bins = 365, color = 'skyblue')
plt.show()
```



Now, lets see if we have any null values in our dataset and how many we have

```
In [ ]: import pandas as pd

# Assuming your DataFrame is named 'airbnb_data'
missing_values = air_df.isnull().sum()

# Calculate the percentage of missing values for each column
total_rows = len(air_df)
missing_percentage = (missing_values / total_rows) * 100

print("Missing values by column:")
print(missing_values)
print("\nPercentage of missing values by column:")
print(missing_percentage)
```

Missing values by column:

id	0
name	16
host_id	0
host_name	21
neighbourhood_group	0
neighbourhood	0
latitude	0
longitude	0
room_type	0
price	0
minimum_nights	0
number_of_reviews	0
last_review	10052
reviews_per_month	10052
calculated_host_listings_count	0
availability_365	0

dtype: int64

Percentage of missing values by column:

id	0.000000
name	0.032723
host_id	0.000000
host_name	0.042949
neighbourhood_group	0.000000
neighbourhood	0.000000
latitude	0.000000
longitude	0.000000
room_type	0.000000
price	0.000000
minimum_nights	0.000000
number_of_reviews	0.000000
last_review	20.558339
reviews_per_month	20.558339
calculated_host_listings_count	0.000000
availability_365	0.000000

dtype: float64

Lets calculate the percentage of missing data of our ENTIRE data set

```
In [ ]: total_entries = air_df.size
missing_values_total = air_df.isnull().sum().sum()

# Calculate the percentage of missing values in the entire dataset
missing_percentage_total = (missing_values_total / total_entries) * 100

print("Total missing values:", missing_values_total)
print("Total entries:", total_entries)
print("Percentage of missing values in the entire dataset:", missing_percent
```

Total missing values: 20141

Total entries: 782320

Percentage of missing values in the entire dataset: 2.574521934758155

Now, from our histograms alone we can see that we do have outliers. Lets check which features have outliers based on IQR

```
In [ ]: # Specify the features you want to check for outliers
features_to_check = ['price', 'minimum_nights', 'number_of_reviews']

for feature in features_to_check:
    # Calculate the first quartile (Q1) and third quartile (Q3)
    Q1 = air_df[feature].quantile(0.25)
    Q3 = air_df[feature].quantile(0.75)

    # Calculate the interquartile range (IQR)
    IQR = Q3 - Q1

    # Determine the outlier step (1.5 times the IQR)
    outlier_step = 1.5 * IQR

    # Identify outliers
    outliers = air_df[(air_df[feature] < (Q1 - outlier_step)) | (air_df[feat

    print("Outliers identified in", feature, ":", outliers)
```

Outliers identified in price :			id
name	host_id	\	
61	15396	Sunny & Spacious Chelsea Apartment	6027
8			
85	19601	perfect for a family or small group	7430
3			
103	23686	2000 SF 3br 2bath West Village private townhouse	9379
0			
114	26933	2 BR / 2 Bath Duplex Apt with patio! East Village	7206
2			
121	27659	3 Story Town House in Park Slope	11958
8			
...	
...			
48758	36420289	Rustic Garden House Apt, 2 stops from Manhattan	7321139
3			
48833	36450896	Brand New 3-Bed Apt in the Best Location of FiDi	2974181
3			
48839	36452721	Massage Spa. Stay overnight. Authors Artist dr...	27407996
4			
48842	36453160	LUXURY MANHATTAN PENTHOUSE+HUDSON RIVER+EMPIRE...	22417137
1			
48856	36457700	Large 3 bed, 2 bath , garden , bbq , all you need	6699339
5			

	host_name	neighbourhood_group	neighbourhood	\
61	Petra	Manhattan	Chelsea	
85	Maggie	Brooklyn	Brooklyn Heights	
103	Ann	Manhattan	West Village	
114	Bruce	Manhattan	East Village	
121	Vero	Brooklyn	South Slope	
...	
48758	LaGabrell	Queens	Long Island City	
48833	Yue	Manhattan	Financial District	
48839	Richard	Brooklyn	Sheepshead Bay	
48842	LuxuryApartmentsByAmber	Manhattan	Chelsea	
48856	Thomas	Brooklyn	Bedford-Stuyvesant	

	latitude	longitude	room_type	price	minimum_nights	\
61	40.74623	-73.99530	Entire home/apt	375	180	
85	40.69723	-73.99268	Entire home/apt	800	1	
103	40.73096	-74.00319	Entire home/apt	500	4	
114	40.72540	-73.98157	Entire home/apt	350	2	
121	40.66499	-73.97925	Entire home/apt	400	2	
...	
48758	40.75508	-73.93258	Entire home/apt	350	2	
48833	40.70605	-74.01042	Entire home/apt	475	2	
48839	40.59866	-73.95661	Private room	800	1	
48842	40.75204	-74.00292	Entire home/apt	350	1	
48856	40.68886	-73.92879	Entire home/apt	345	4	

	number_of_reviews	last_review	reviews_per_month	\
61	5	2018-11-03	0.12	
85	25	2016-08-04	0.24	
103	46	2019-05-18	0.55	
114	7	2017-08-09	0.06	

121	16	2018-12-30	0.24
...
48758	0	NaN	NaN
48833	0	NaN	NaN
48839	0	NaN	NaN
48842	0	NaN	NaN
48856	0	NaN	NaN

	calculated_host_listings_count	availability_365
61	1	180
85	1	7
103	2	243
114	4	298
121	2	216
...
48758	1	364
48833	1	64
48839	1	23
48842	1	9
48856	3	354

[2972 rows x 16 columns]

Outliers identified in minimum_nights :

	id
name host_id \	
6 5121	BlissArtsSpace! 735
6	
14 6090	West Village Nest – Superhost 1197
5	
29 9657	Modern 1 BR / NYC / EAST VILLAGE 2190
4	
36 11452	Clean and Quiet in Brooklyn 735
5	
45 12627	Entire apartment in central Brooklyn neighborh... 4967
0	
...	...
...	
48810 36445121	UWS Spacious Master Bedroom Sublet 27401445
3	
48843 36453642	☆ HUGE, SUNLIT Room – 3 min walk from Train ! 5396611
5	
48871 36475746	A LARGE ROOM – 1 MONTH MINIMUM – WASHER&DRYER 14400870
1	
48879 36480292	Gorgeous 1.5 Bdr with a private yard– Williams... 54033
5	
48882 36482231	Bushwick _ Myrtle–Wyckoff 6605889
6	

	host_name	neighbourhood_group	neighbourhood	latitude \
6	Garon	Brooklyn	Bedford–Stuyvesant	40.68688
14	Alina	Manhattan	West Village	40.73530
29	Dana	Manhattan	East Village	40.72920
36	Vt	Brooklyn	Bedford–Stuyvesant	40.68876
45	Rana	Brooklyn	Prospect–Lefferts Gardens	40.65944
...
48810	Dagmara	Manhattan	Upper West Side	40.79952
48843	Nora	Brooklyn	Bedford–Stuyvesant	40.69635

48871	Ozzy Ciao	Manhattan	Harlem	40.82233
48879	Lee	Brooklyn	Williamsburg	40.71728
48882	Luisa	Brooklyn	Bushwick	40.69652

	longitude	room_type	price	minimum_nights	number_of_reviews
\					
6	-73.95596	Private room	60	45	49
14	-74.00525	Entire home/apt	120	90	27
29	-73.98542	Entire home/apt	180	14	29
36	-73.94312	Private room	35	60	0
45	-73.96238	Entire home/apt	150	29	11
...
48810	-73.96003	Private room	75	30	0
48843	-73.93743	Private room	45	29	0
48871	-73.94687	Private room	35	29	0
48879	-73.94394	Entire home/apt	120	20	0
48882	-73.91079	Private room	40	20	0

	last_review	reviews_per_month	calculated_host_listings_count	\
6	2017-10-05	0.40	1	
14	2018-10-31	0.22	1	
29	2019-04-19	0.24	1	
36	NaN	NaN	1	
45	2019-06-05	0.49	1	
...	
48810	NaN	NaN	1	
48843	NaN	NaN	2	
48871	NaN	NaN	2	
48879	NaN	NaN	1	
48882	NaN	NaN	1	

	availability_365
6	0
14	0
29	67
36	365
45	95
...	...
48810	90
48843	341
48871	31
48879	22
48882	31

[6607 rows x 16 columns]

Outliers identified in number_of_reviews :

name	host_id	\	id
3	3831	Cozy Entire Floor of Brownstone	486
9			
5	5099	Large Cozy 1 BR Apartment In Midtown East	732
2			
7	5178	Large Furnished Room Near B'way	896
7			
8	5203	Cozy Clean Guest Room – Family Apt	749
0			
9	5238	Cute & Cozy Lower East Side 1 bdrm	754

```

9
...      ...
...
40104  31123611      JFK Airport Great place to stay 6 minutes away  23225188
1
40297  31249784      Studio Apartment 6 minutes from JFK Airport  23225188
1
40424  31336245      Jfk crash pad 1-2persons in SHARED space  23225188
1
42075  32678719  Enjoy great views of the City in our Deluxe Room!  24436158
9
42076  32678720      Great Room in the heart of Times Square!  24436158
9

```

```

      host_name neighbourhood_group      neighbourhood  latitude  longitud
e \
3      LisaRoxanne      Brooklyn      Clinton Hill  40.68514  -73.9597
6
5      Chris      Manhattan      Murray Hill  40.74767  -73.9750
0
7      Shunichi      Manhattan      Hell's Kitchen  40.76489  -73.9849
3
8      MaryEllen      Manhattan      Upper West Side  40.80178  -73.9672
3
9      Ben      Manhattan      Chinatown  40.71344  -73.9903
7

```

```

...      ...      ...      ...      ...
...
40104      Lakshmee      Queens      Jamaica  40.66823  -73.7837
4
40297      Lakshmee      Queens      Jamaica  40.66793  -73.7845
2
40424      Lakshmee      Queens      Jamaica  40.66715  -73.7834
6
42075      Row NYC      Manhattan      Theater District  40.75918  -73.9880
1
42076      Row NYC      Manhattan      Theater District  40.75828  -73.9887
6

```

```

      room_type  price  minimum_nights  number_of_reviews  last_review
\
3      Entire home/apt      89      1      270  2019-07-05
5      Entire home/apt      200      3      74  2019-06-22
7      Private room      79      2      430  2019-06-24
8      Private room      79      2      118  2017-07-21
9      Entire home/apt      150      1      160  2019-06-09
...      ...      ...      ...      ...
40104      Shared room      40      1      65  2019-07-06
40297      Private room      67      1      95  2019-07-05
40424      Shared room      39      1      65  2019-07-07
42075      Private room      100      1      156  2019-07-07
42076      Private room      199      1      82  2019-07-07

```

```

      reviews_per_month  calculated_host_listings_count  availability_365
3      4.64      1      194
5      0.59      1      129

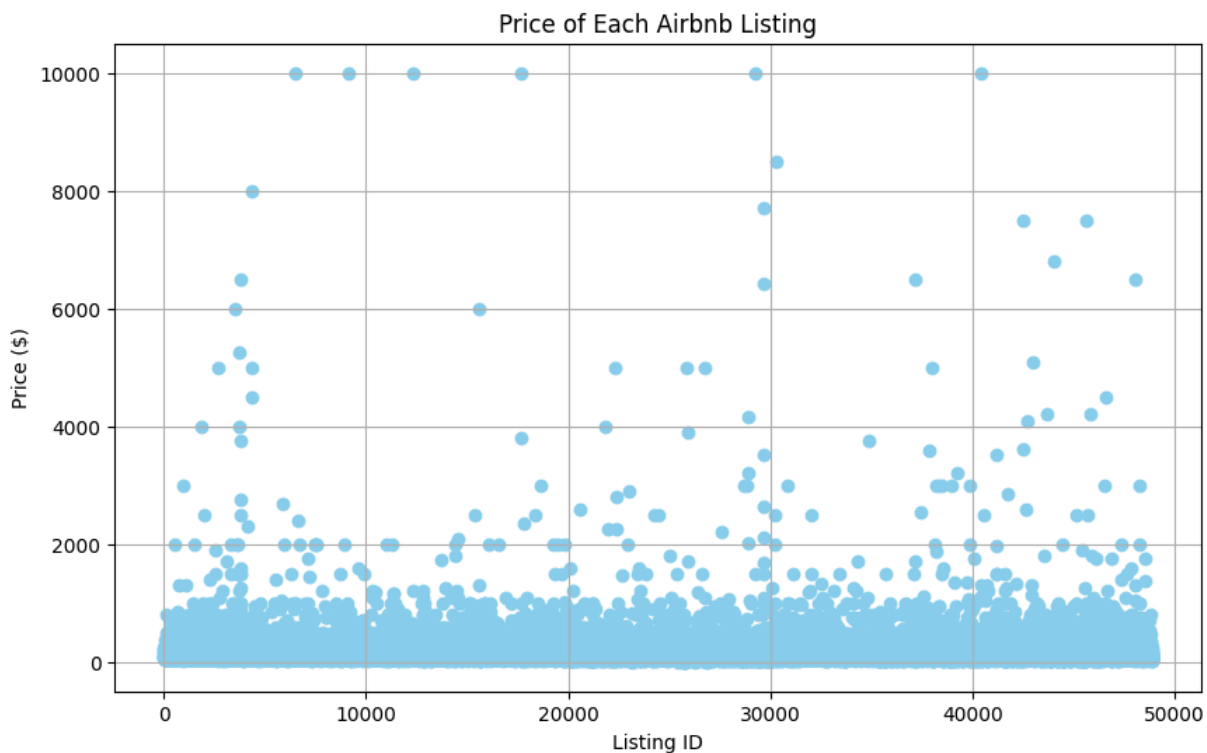
```

7	3.47	1	220
8	0.99	1	0
9	1.33	4	188
...
40104	10.00	8	346
40297	15.32	8	145
40424	10.60	8	320
42075	58.50	9	299
42076	27.95	9	299

[6021 rows x 16 columns]

Let's try and visualize some of the outliers using a scatterplot of the prices of all the listings

```
In [ ]: # Plotting the price of each Airbnb listing
plt.figure(figsize=(10, 6))
plt.scatter(range(len(air_df)), air_df['price'], color='skyblue')
plt.title('Price of Each Airbnb Listing')
plt.xlabel('Listing ID')
plt.ylabel('Price ($)')
plt.grid(True)
plt.show()
```



Now, let's look for correlations. We are going to do this by plotting the correlation matrix for all features in our data

```
In [ ]: columns_except_price = air_df.drop(columns=['price', 'id', 'host_name', 'name'])
X = pd.get_dummies(columns_except_price, drop_first=True)
X.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12 entries, 0 to 11
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   calculated_host_listings_count        12 non-null     uint8
1   host_id                               12 non-null     uint8
2   last_review                           12 non-null     uint8
3   latitude                              12 non-null     uint8
4   longitude                             12 non-null     uint8
5   minimum_nights                        12 non-null     uint8
6   neighbourhood                         12 non-null     uint8
7   neighbourhood_group                   12 non-null     uint8
8   number_of_reviews                    12 non-null     uint8
9   reviews_per_month                    12 non-null     uint8
10  room_type                             12 non-null     uint8
dtypes: uint8(11)
memory usage: 264.0 bytes
```

```
In [ ]: # Create the target (or outcome) field
        y = air_df['price']
```

```
In [ ]: corr = pd.concat([X, y], axis=1).corr()

# Include information about values
fig, ax = plt.subplots(ncols=1, rows=1, figsize=(11, 7))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="RdBu", center=0, ax=ax);
ax.set_title('NY Airbnb Feature Correlation Matrix')
ax.set_xlabel('Airbnb Features')
ax.set_ylabel('Airbnb Features')

plt.tight_layout()
plt.show()
```



Now, let's examine some of the correlations more closely using histograms, box plots, and scatter plots

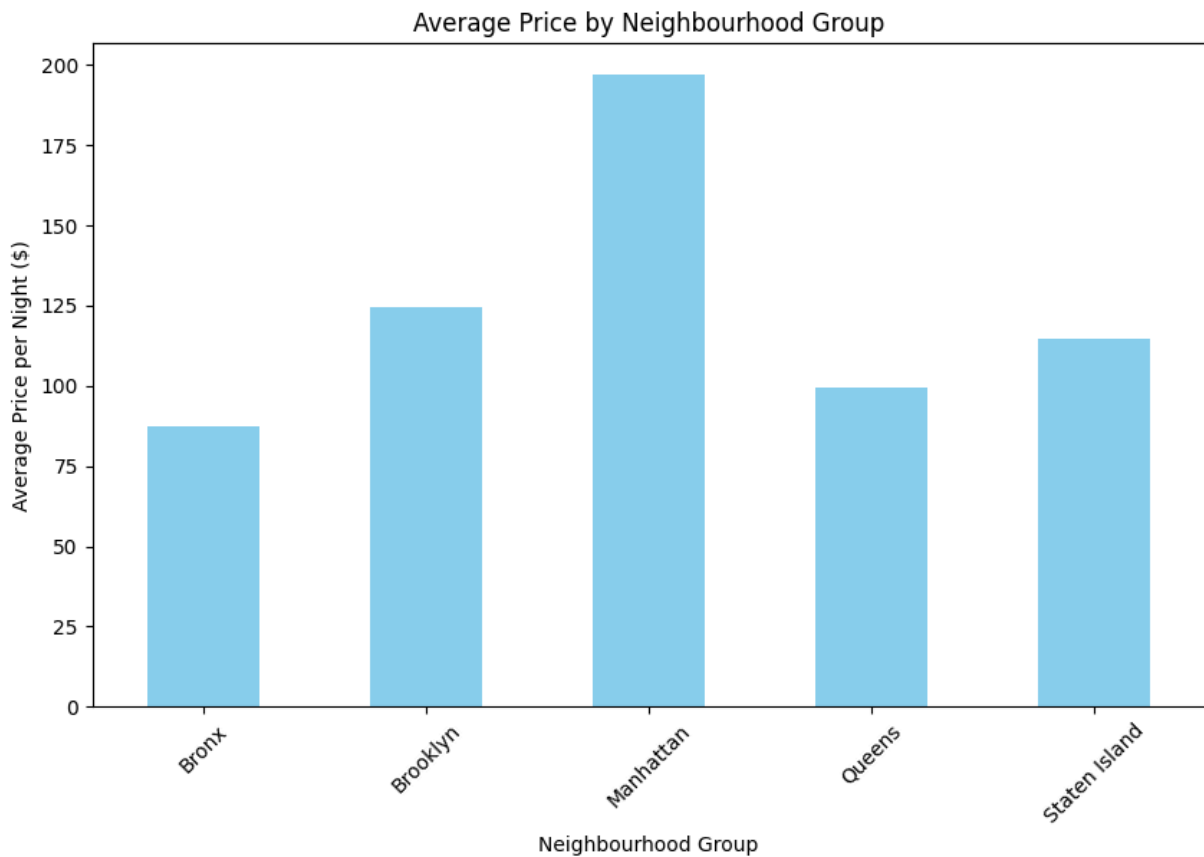
```
In [ ]: # Group data by 'neighbourhood_group' and calculate mean
average_price_by_neighbourhood_group = air_df.groupby('neighbourhood_group')

# Plot the average price by neighbourhood_group
average_price_by_neighbourhood_group.plot(kind='bar', figsize=(10, 6), color

# Set the title and labels
plt.title('Average Price by Neighbourhood Group')
plt.xlabel('Neighbourhood Group')
plt.ylabel('Average Price per Night ($)')

# Rotate x axis labels by 45 degrees
plt.xticks(rotation=45)

# Show plot
plt.show()
```



Lets plot the Histogram for the top 15 neighborhoods


```
In [ ]: # Get the counts of each neighbourhood
neighbourhood_counts = air_df['neighbourhood'].value_counts()

# Select the top 15 most common neighbourhoods
top_neighbourhoods = neighbourhood_counts.head(15)

# Plot the distribution of neighbourhood
plt.figure(figsize=(10, 6))
top_neighbourhoods.plot(kind='bar', color='skyblue')
plt.title('Distribution of Top 15 Neighbourhoods by Frequency')
plt.xlabel('Neighbourhood')
plt.ylabel('Frequency')

# Rotate x axis labels by 45 degrees
plt.xticks(rotation=45)
plt.tight_layout()

plt.show()
```



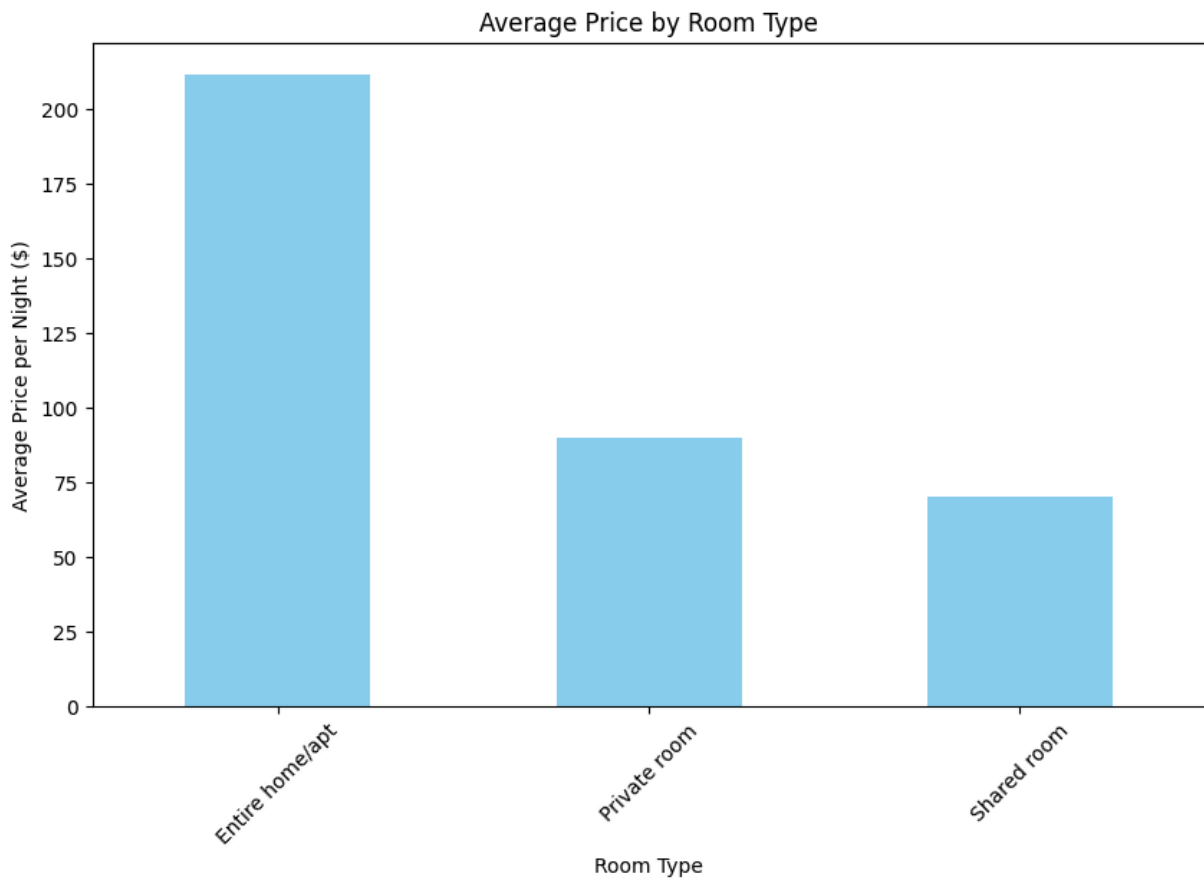
```
In [ ]: # Group data by 'room_type' and calculate mean
average_price_by_room_type = air_df.groupby('room_type')['price'].mean()

# Plot the average price by room_type
average_price_by_room_type.plot(kind='bar', figsize=(10, 6), color='skyblue')

# Set the title and labels
plt.title('Average Price by Room Type')
plt.xlabel('Room Type')
plt.ylabel('Average Price per Night ($)')

# Rotate x axis labels by 45 degrees
plt.xticks(rotation=45)

# Show plot
plt.show()
```



```
In [ ]: # Get the top 15 most common neighbourhoods
top_neighbourhoods = air_df['neighbourhood'].value_counts().head(15).index.t


# Filter the DataFrame to include only data for the top 15 neighbourhoods
top_neighbourhood_data = air_df[air_df['neighbourhood'].isin(top_neighbourhc

# Group the data by 'neighbourhood' and calculate the mean price within each
average_price_by_neighbourhood = top_neighbourhood_data.groupby('neighbourhc

# Plot the average price for the top 15 neighbourhoods
plt.figure(figsize=(10, 6))
average_price_by_neighbourhood.plot(kind='bar', color='skyblue')
plt.title('Average Price for Top 15 Neighbourhoods by Price')
plt.xlabel('Neighbourhood')
plt.ylabel('Average Price per Night ($)')

# Rotate x axis labels by 45 degrees
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

```
In [ ]: # Plot scatter plot of price vs availability_365
plt.figure(figsize=(10, 6))

plt.scatter(air_df['availability_365'], air_df['price'], alpha=0.5, color='s')

plt.title('Price vs Availability (365 days)')
plt.xlabel('Availability (365 days)')
plt.ylabel('Price per Night ($)')
plt.grid(True)

plt.show()
```

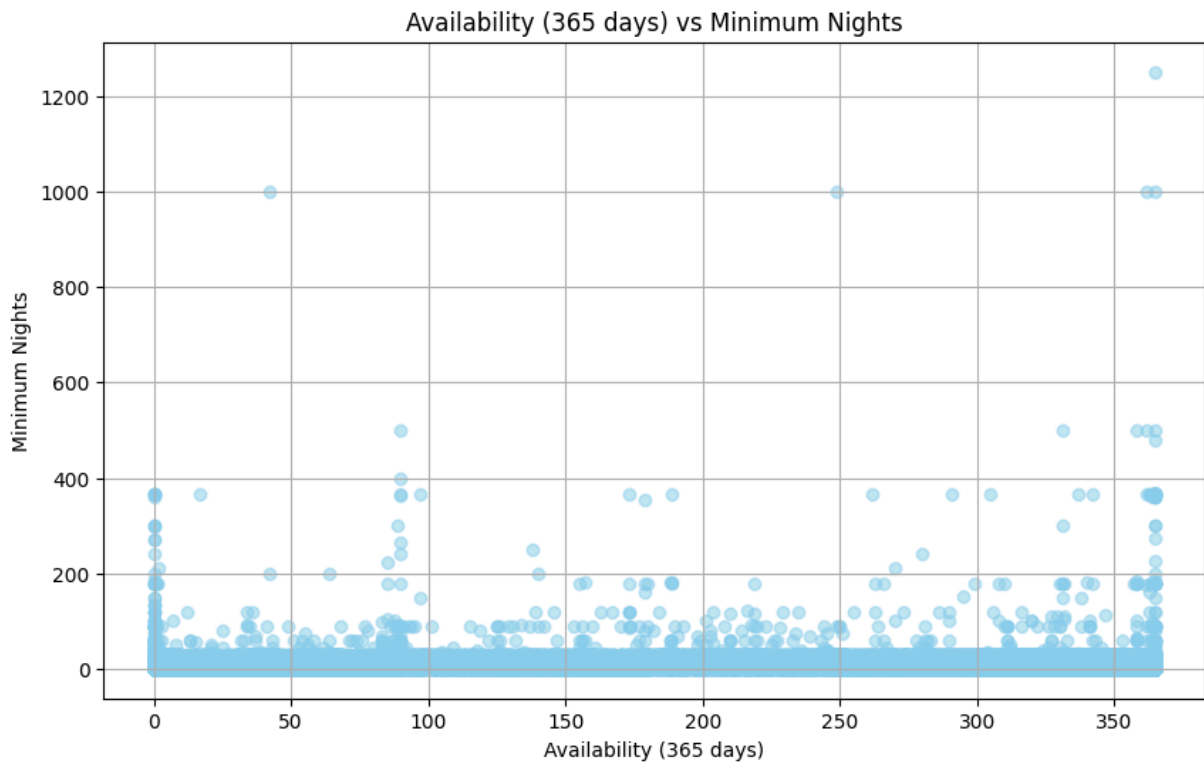


```
In [ ]: # Plot scatter plot of availability_365 vs minimum_nights
plt.figure(figsize=(10, 6))

plt.scatter(air_df['availability_365'], air_df['minimum_nights'], alpha=0.5,

plt.title('Availability (365 days) vs Minimum Nights')
plt.xlabel('Availability (365 days)')
plt.ylabel('Minimum Nights')
plt.grid(True)

plt.show()
```

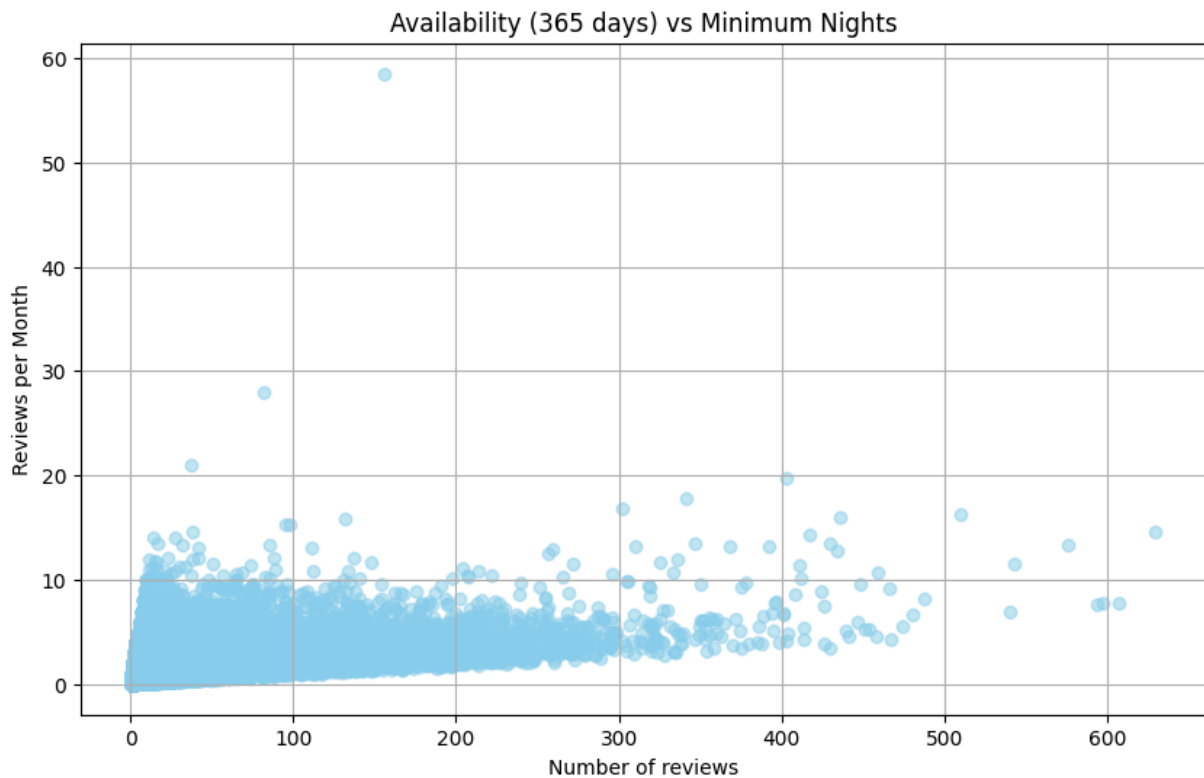


```
In [ ]: # Plot scatter plot of availability_365 vs minimum_nights
plt.figure(figsize=(10, 6))

plt.scatter(air_df['number_of_reviews'], air_df['reviews_per_month'], alpha=

plt.title('Availability (365 days) vs Minimum Nights')
plt.xlabel('Number of reviews')
plt.ylabel('Reviews per Month')
plt.grid(True)

plt.show()
```



```
In [ ]: def report(df):  
    col = []  
    d_type = []  
    uniques = []  
    n_uniques = []  
  
    for i in df.columns:  
        col.append(i)  
        d_type.append(df[i].dtypes)  
        uniques.append(df[i].unique()[:5])  
        n_uniques.append(df[i].nunique())  
  
    return pd.DataFrame({'Column': col, 'd_type': d_type, 'unique_sample': u
```

```
In [ ]: report(air_df)
```

Out []:

	Column	d_type	unique_sample	n_uniques
0	id	int64	[2539, 2595, 3647, 3831, 5022]	48895
1	name	object	[Clean & quiet apt home by the park, Skylit Mi...	47905
2	host_id	int64	[2787, 2845, 4632, 4869, 7192]	37457
3	host_name	object	[John, Jennifer, Elisabeth, LisaRoxanne, Laura]	11452
4	neighbourhood_group	object	[Brooklyn, Manhattan, Queens, Staten Island, B...	5
5	neighbourhood	object	[Kensington, Midtown, Harlem, Clinton Hill, Ea...	221
6	latitude	float64	[40.64749, 40.75362, 40.80902, 40.68514, 40.79...	19048
7	longitude	float64	[-73.97237, -73.98377, -73.9419, -73.95976, -7...	14718
8	room_type	object	[Private room, Entire home/apt, Shared room]	3
9	price	int64	[149, 225, 150, 89, 80]	674
10	minimum_nights	int64	[1, 3, 10, 45, 2]	109
11	number_of_reviews	int64	[9, 45, 0, 270, 74]	394
12	last_review	object	[2018-10-19, 2019-05-21, nan, 2019-07-05, 2018...	1764
13	reviews_per_month	float64	[0.21, 0.38, nan, 4.64, 0.1]	937
14	calculated_host_listings_count	int64	[6, 2, 1, 4, 3]	47
15	availability_365	int64	[365, 355, 194, 0, 129]	366

Lets change some object types to Categorical as it will make later analysis easier.

More specifically, we are going to change neighbourhood_group and room_type to categorical

```
In [ ]: air_df['neighbourhood_group'] = air_df['neighbourhood_group'].astype('category')
air_df['room_type'] = air_df['room_type'].astype('category')
```

```
In [ ]: air_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     48895 non-null  int64
1   name                                  48879 non-null  object
2   host_id                               48895 non-null  int64
3   host_name                             48874 non-null  object
4   neighbourhood_group                   48895 non-null  category
5   neighbourhood                         48895 non-null  object
6   latitude                             48895 non-null  float64
7   longitude                             48895 non-null  float64
8   room_type                             48895 non-null  category
9   price                                 48895 non-null  int64
10  minimum_nights                        48895 non-null  int64
11  number_of_reviews                     48895 non-null  int64
12  last_review                           38843 non-null  object
13  reviews_per_month                     38843 non-null  float64
14  calculated_host_listings_count         48895 non-null  int64
15  availability_365                       48895 non-null  int64
dtypes: category(2), float64(3), int64(7), object(4)
memory usage: 5.3+ MB

```

We know that we have null values for 1052 entries, lets fill these null values with the mean of the column

```

In [ ]: air_df.fillna(air_df.mean(), inplace=True)
air_df.info()
air_df.head()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     48895 non-null  int64
1   name                                  48879 non-null  object
2   host_id                               48895 non-null  int64
3   host_name                             48874 non-null  object
4   neighbourhood_group                   48895 non-null  category
5   neighbourhood                         48895 non-null  object
6   latitude                             48895 non-null  float64
7   longitude                             48895 non-null  float64
8   room_type                             48895 non-null  category
9   price                                 48895 non-null  int64
10  minimum_nights                        48895 non-null  int64
11  number_of_reviews                     48895 non-null  int64
12  last_review                           38843 non-null  object
13  reviews_per_month                     48895 non-null  float64
14  calculated_host_listings_count         48895 non-null  int64
15  availability_365                       48895 non-null  int64
dtypes: category(2), float64(3), int64(7), object(4)
memory usage: 5.3+ MB

```

Out []:	id	name	host_id	host_name	neighbourhood_group	neighbourhood
0	2539	Clean & quiet apt home by the park	2787	John	Brooklyn	Kensington
1	2595	Skylit Midtown Castle	2845	Jennifer	Manhattan	Midtown
2	3647	THE VILLAGE OF HARLEM....NEW YORK !	4632	Elisabeth	Manhattan	Harlem
3	3831	Cozy Entire Floor of Brownstone	4869	LisaRoxanne	Brooklyn	Clinton Hill
4	5022	Entire Apt: Spacious Studio/Loft by central park	7192	Laura	Manhattan	East Harlem

Now, let's develop our first model, a Decision Tree Regression model

```
In [ ]: # Based on the correlations, set predictors and outcome
predictors = ['neighbourhood_group', 'latitude', 'longitude', 'room_type', '
outcome = 'price'

# Create dummy variables for categorical predictors
X = pd.get_dummies(air_df[predictors], drop_first=True)
y = air_df[outcome]

# select the columns that are objects and typecast them to category
cat_cols = X.select_dtypes(include='object').columns
X[cat_cols] = X[cat_cols].astype('category')
```

Create our training and testing data:

```
In [ ]: # partition 80% of the data into training and 20% into testing sets
train_X, test_X, train_y, test_y = train_test_split(X, y, train_size=0.8, ra
```

Lets fine-tune the hyperparameters of the model using GridSearchCV.

```
In [ ]: # Define the Decision Tree Regression model
dt_model = DecisionTreeRegressor()

# Define the parameter grid for GridSearchCV
param_grid = {
    'max_depth': [3, 5, 7, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 3, 5, 7, 10]
}

# Perform GridSearchCV
```

```

grid_search = GridSearchCV(estimator=dt_model, param_grid=param_grid, cv=5)
grid_search.fit(train_X, train_y)

# Print the best parameters
print("Best parameters found by GridSearchCV:")
print(grid_search.best_params_)

# Make predictions with the best model
best_model = grid_search.best_estimator_
test_pred_y = best_model.predict(test_X)

# Calculate mean absolute error on the test set
mae = mean_absolute_error(test_y, test_pred_y)
print("Mean Squared Error on Test Set:", mae)

# Compare with LassoCV model
lasso_model = LassoCV(cv=5)
lasso_model.fit(train_X, train_y)
lasso_test_pred_y = lasso_model.predict(test_X)

lasso_mae = mean_absolute_error(test_y, lasso_test_pred_y)
print("Mean Squared Error with LassoCV:", lasso_mae)

```

Best parameters found by GridSearchCV:
{'max_depth': 7, 'min_samples_leaf': 10, 'min_samples_split': 2}
Mean Squared Error on Test Set: 69.73328037182786
Mean Squared Error with LassoCV: 74.92954822821314

Now, lets create the actual model with the best hyperparameters we found in the previous step.

```

In [ ]: # Initialize the DecisionTreeRegressor model with the best parameters
best_dt_model = DecisionTreeRegressor(criterion="squared_error", max_depth=5)

# Fit the model to the training data
best_dt_model.fit(train_X, train_y)

```

```

Out[ ]: ▼ DecisionTreeRegressor ⓘ ?
DecisionTreeRegressor(max_depth=5, min_samples_leaf=3)

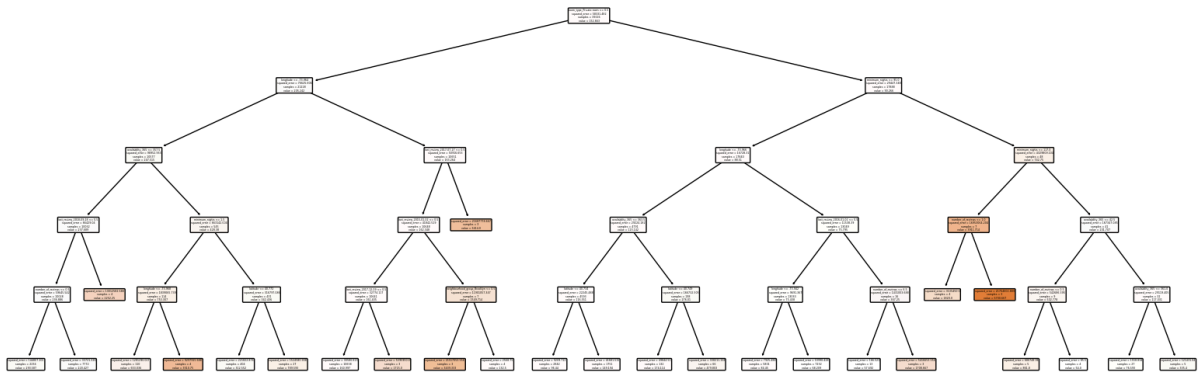
```

Lets see the tree to verify the results.

```

In [ ]: plt.figure(figsize=(20, 7))
plot_tree(best_dt_model, feature_names=train_X.columns.to_list(), filled = True)
plt.show()

```

Now, let's begin quickly testing the model. First, let's use MAE followed by cross validation with MAPE as the scoring.

```
In [ ]: # Make predictions on the validation data using the best DT Regressor model
val_pred_y = best_dt_model.predict(test_X)

# Calculate the MAE to quickly evaluate performance
mae = mean_absolute_error(test_y, val_pred_y)

print("Mean Absolute Percentage Error on Validation Data:", mae)
```

Mean Absolute Percentage Error on Validation Data: 74.35689405333234

```
In [ ]: # Perform cross-validation to evaluate the model's performance
cv_scores = cross_val_score(best_dt_model, X, y, cv=5, scoring='neg_mean_abs

# Take the negative mean of the cv scores to get MAE
cv_mae = -cv_scores.mean()

print("Cross-Validation Mean Absolute Error:", cv_mae)
```

Cross-Validation Mean Absolute Error: 73.88596354290772

Let's implement a Gradient Boosting Regression model! This builds smaller DT's off of the previous ones and should account for small errors within them, thus increasing performance.

```
In [ ]: # Create the Gradient Boosting Regression model
gb_model = GradientBoostingRegressor()

# Fit the model to the training data
gb_model.fit(train_X, train_y)

# Make predictions on the test data
y_pred = gb_model.predict(test_X)

# Quickly evaluate the model
mae = mean_absolute_error(test_y, y_pred)
print("Mean Absolute Error (MAE):", mae)
```

Mean Absolute Error (MAE): 69.45034711502922


Since we can't see all the trees that the model creates, let's try and visualize the performance using a scatter plot of the actual and predicted values using the Gradient Boosting Model.

```
In [ ]: # Save the predicted values
predicted_values = gb_model.predict(test_X)

# Plotting the actual prices
plt.figure(figsize=(10, 6))
plt.scatter(test_X.index, test_y, color='skyblue', label='Actual Price')

# Plotting the predicted prices
plt.scatter(test_X.index, predicted_values, color='orange', label='Predicted Price')

plt.title('Actual vs Predicted Prices with the Gradient Boosting Model')
plt.xlabel('Listing ID')
plt.ylabel('Price ($)')
plt.legend()
plt.grid(True)
plt.show()
```



Now, let's compare the accuracy to a dummy model. For the sake of simplicity and understanding, let's use the mean dummy model.

```
In [ ]: # Define the Dummy model using the mean value as the strategy (predict the mean)
dummy_model = DummyRegressor(strategy='mean')

# Train the Dummy model
dummy_model.fit(train_X, train_y)
```

```

# Make predictions with the Dummy model
dummy_test_pred_y = dummy_model.predict(test_X)

# Calculate mean absolute error for the Dummy model
dummy_mae = mean_absolute_error(test_y, dummy_test_pred_y)
print("Mean Absolute Error with Dummy model (mean strategy):", dummy_mae)

```

Mean Absolute Error with Dummy model (mean strategy): 93.34945129158264

Now, let's compare accuracy and error metrics for each model (Including LassoCV so we can compare it to the Decision Tree Regression model)

```

In [ ]: # Define a function to calculate multiple evaluation metrics
def calculate_metrics(true_y, pred_y):
    mae = mean_absolute_error(true_y, pred_y)
    mse = mean_squared_error(true_y, pred_y)
    rmse = np.sqrt(mse)
    r2 = r2_score(true_y, pred_y)
    n = len(true_y)
    p = test_X.shape[1]
    adj_r2 = 1 - (1 - r2) * ((n - 1) / (n - p - 1)) # Compute Adjusted R-squared
    return mae, mse, rmse, r2, adj_r2

# Evaluate Decision Tree model
dt_mae, dt_mse, dt_rmse, dt_r2, dt_adj_r2 = calculate_metrics(test_y, test_pred_y)

# Evaluate LassoCV model
lasso_mae, lasso_mse, lasso_rmse, lasso_r2, lasso_adj_r2 = calculate_metrics(test_y, test_pred_y)

# Evaluate Gradient Boosting model
gb_mae, gb_mse, gb_rmse, gb_r2, gb_adj_r2 = calculate_metrics(test_y, y_pred)

# Evaluate Dummy model
dummy_mae, dummy_mse, dummy_rmse, dummy_r2, dummy_adj_r2 = calculate_metrics(test_y, test_pred_y)

# Print the evaluation metrics for Decision Tree, Gradient Boosting, and Dummy model
print("Decision Tree Model:")
print("MAE:", dt_mae)
print("MSE:", dt_mse)
print("RMSE:", dt_rmse)
print("R-squared:", dt_r2)
print("Adjusted R-squared:", dt_adj_r2)
print()

print("LassoCV Model:")
print("MAE:", lasso_mae)
print("MSE:", lasso_mse)
print("RMSE:", lasso_rmse)
print("R-squared:", lasso_r2)
print("Adjusted R-squared:", lasso_adj_r2)
print()

print("Gradient Boosting Model:")
print("MAE:", gb_mae)
print("MSE:", gb_mse)
print("RMSE:", gb_rmse)

```

```

print("R-squared:", gb_r2)
print("Adjusted R-squared:", gb_adj_r2)
print()

print("Dummy Model:")
print("MAE:", dummy_mae)
print("MSE:", dummy_mse)
print("RMSE:", dummy_rmse)
print("R-squared:", dummy_r2)
print("Adjusted R-squared:", dummy_adj_r2)
print()

```

Decision Tree Model:

MAE: 69.73328037182786
MSE: 49996.24858786161
RMSE: 223.59840918007805
R-squared: 0.09160877191822858
Adjusted R-squared: -0.10972631536526234

LassoCV Model:

MAE: 74.92954822821314
MSE: 49722.400746358224
RMSE: 222.98520297624734
R-squared: 0.09658436476923304
Adjusted R-squared: -0.10364793619270851

Gradient Boosting Model:

MAE: 69.45034711502922
MSE: 53073.18201817909
RMSE: 230.376175022894
R-squared: 0.03570339068587447
Adjusted R-squared: -0.1780225194744527

Dummy Model:

MAE: 93.34945129158264
MSE: 55038.31691017777
RMSE: 234.60246569500876
R-squared: -1.5141481737313e-06
Adjusted R-squared: -0.2216410301525289

Gradient Boosting performed the best with the basic DT Regressor not far behind. LassoCV did a pretty good job for the ease of implementation. The Dummy Model performed drastically worse than any of the other models as expected.